

Cloud-Native Renewable Energy Marketplace: A Comprehensive Peer-to-Peer Trading Platform Using Google Cloud Technologies

K. Krishnanunni¹, Krishnapriya J.D², Malavika Ajith³,

Vidya Krishnan S⁴, *Divya Udayan J⁵

^{1,2,3,4,5}Department of Computer Science and Engineering,

Amrita School of Computing, Amrita Vishwa Vidyapeetham, Amritapuri, India

Contributing authors:

am.en.u4cse22231@am.students.amrita.edu,

am.en.u4cse22235@am.students.amrita.edu, am.en.u4cse22242@am.students.amrita.edu,

am.en.u4cse22360@am.students.amrita.edu

*Corresponding author: divyaudayanj@am.amrita.edu

Abstract—This work presents a cloud-native renewable energy marketplace platform that employs a digital credit system to enable decentralized peer-to-peer (P2P) trading. The platform, developed on the Google Cloud Platform with Firebase, BigQuery, Dialogflow, and Stripe, remedies major inefficiencies in traditional energy markets. A six-month pilot rollout in rural Kerala, India, outperformed blockchain peers with sub-150ms transaction processing times, 99.9% availability, and 96.4% cost reduction. The platform benefited 5,138 users, achieving 94.2% satisfaction (n=487, 95% CI: 91.8%–96.6%), meeting the specific needs of rural communities, where widespread distributed renewable installations are present but the technical infrastructure is limited. It facilitated 12,847 energy transactions worth Rs.2.3 million. The main innovations are regulatory-compliant energy credit authentication, smart buyer-seller matching with a success rate of 97.3% and AI-driven dynamic pricing with a performance improvement. 15%. The system produced a demonstrable economic impact in which consumers received an average of Rs.18,400 per year, targeting the Sustainable Development Goals (SDGs) of the United Nations for clean energy and rural upliftment. By facilitating low-cost, inclusive energy trading and reducing negative carbon emissions, the platform creates a pathway towards long term environmental, economic and social sustainability.

Index Terms—Renewable energy trading, cloud computing, peer-to-peer energy markets, Google Cloud Platform, artificial intelligence and smart grid.

I. INTRODUCTION

The growth of decentralized renewable energy systems has brought unprecedented prospects for decentralized power generation, and worldwide distributed solar capacity is predicted to total 146 GW in 2024 [1], [2]. Yet, small power producers from renewable energy sources struggle with enormous barriers in trading surplus energy due to poor market mechanisms, overcomplicated rules, and limited technology [3], [4]. Old school centralized energy markets are unable to fit into the volatile, distributed nature of renewable energy sources, leading to economic inefficiency and low incentives for clean energy adoption [5].

Current peer-to-peer (P2P) energy trading systems are plagued with various fundamental drawbacks. Blockchain systems have high computational burden and latency (average 3–15 seconds per transaction) [6], whereas existing cloud platforms do not have end-to-end integration of payment processing, regulation, and user-friendly interfaces. Most importantly, most of these solutions do not satisfy the particular requirements of rural communities, where there are widespread distributed renewable installations but limited technical infrastructure [7].

To counter such challenges, this paper introduces a holistic cloud-native Renewable Energy Marketplace Platform based on Google Cloud technologies to facilitate effective P2P trading of renewable energy credits. While physical trading of energy entails intricate grid infrastructure, our system centers on standardized renewable energy certificates (RECs) for authenticated clean energy generation, and hence presents a viable and scalable means of monetization of clean energy [8], [9].

Novelty and Contribution of this paper are as follows: This paper describes the very first cloud-native, peer-to-peer renewable energy trading platform developed and validated at scale using Google Cloud services. In contrast to solutions that would typically rely on blockchain or centralized architectures, our system integrates AI-driven dynamic pricing, real-time renewable credit verification, multilingual progressive web accessibility, and automated regulatory compliance — all in a low-latency, serverless cloud infrastructure. The platform is optimised for rural deployment and has been tested through a six-month pilot, which proved measurable economic, environmental, and social impact.

The proposed research will design and test a cloud-native platform of energy trading specifically targeting rural Indian societies but prioritizing regulatory compliance and technological affordability. Its main aim is to connect the small

scale renewable energy producers to the consumers using an effective market place. The research questions discussed in the study are the following: (1) do cloud-native architectures allow outperforming blockchain-based solutions in energy trading; (2) how can AI-driven algorithms be applied to optimize pricing and matching of renewable markets; and (3) what quantifiable economic value P2P trading can bring to rural communities that have distributed renewable resources.

This work has a few major contributions to the research on decentralized energy systems. To start with, we are offering a complete cloud-native system based on Google Cloud services to end-to-end energy credit trading. Second, we create AI-driven dynamic types of pricing algorithms, which react to market conditions, supply-demand interactions and weather conditions. Third, we introduce a regulatory-compliant credit system that is in line with international standards (I-REC, Green-e), and it is tested by six-month pilot implementation with more than 5,000 users. The platform is more efficient than blockchain-based systems and allows real economic advantages to rural participants. Taken together, such developments contribute to some of the priority UN Sustainable Development Goals, namely, affordable clean energy (SDG 7), economic growth (SDG 8), and climate action (SDG 13).

II. LITERATURE REVIEW AND RELATED WORK

A. Literature Review Methodology

This review utilized a systematic, multi-phase process to find and review previous work on decentralized renewable energy trading systems. The following databases were searched: IEEE Xplore, SpringerLink, ScienceDirect, and Google Scholar.

Search Keywords: “P2P energy trading,” “blockchain energy markets,” “cloud-native energy platforms,” “AI in smart grid,” and “renewable energy in rural areas.”

Inclusion Criteria:

- Peer-reviewed publications from 2015 to 2024.
- Focus on decentralized or peer-to-peer renewable energy markets.
- Studies addressing blockchain or cloud-native energy trading architectures.
- Use of AI or machine learning for pricing, prediction, or matching.
- Relevance to rural deployment, regulation, and real-world validation.

Of the 57 articles claimed, after the removal of duplicates and review of relevance, 23 were included. These were thematically coded as follows.

B. Blockchain-Based Energy Platforms

Blockchain is being studied for decentralized energy markets because of its transparency and decentralization. Such as smart contracts and consensus mechanisms, which create trust by Paudel et al. [3] and Tushar et al. [4]. However, many performance constraints prevent real-time deployment.

Blockchain Advantages:

TABLE I
PERFORMANCE COMPARISON OF P2P ENERGY TRADING PLATFORMS

Platform	Technology	Avg. Time	TPS	Cost	Focus
PowerLedger	Blockchain	8–12s	15	\$2–5	Commercial
Grid+	Blockchain	5–8s	25	\$1–3	Residential
LO3 Energy	Blockchain	10–15s	12	\$3–7	Microgrid
WePower	Blockchain	6–10s	20	\$2–4	Utility
Our Platform	Cloud-Native	<150ms	1000+	\$0.05–0.10	Rural

- Immutable audit trails and distributed consensus for trustless transactions.
- Smart contracts for automated settlements and payment enforcement.

Key Limitations: According to PowerLedger’s reports [10] and Ethereum-based benchmarks [11]:

- Low throughput (15–25 TPS).
- High confirmation latency (8–12 seconds).
- Energy-intensive processing (150–200 kWh per transaction) [12].
- Suboptimal resource utilization in small and medium-scale cloud environments, leading to occasional over-provisioning or under-provisioning of virtual machines [13].

C. Cloud-Native Architectures: A Viable Alternative

Appropriately, cloud-native platforms present significant advantages in latency, scalability, and integration with modern APIs. Kim et al. [11] and Wang et al. [14] provide examples of how cloud-first approaches yield responsive design, scalable deployment, and fewer overhead costs for rural applications.

Our platform takes advantage of these abilities to achieve less than 150ms latency, over 1000 TPS, and compliance without a blockchain. A summary of platform comparisons is presented in Table I.

D. AI in Pricing and Matching Optimization

Various artificial intelligence techniques have been suggested for pricing and electricity demand forecasting. Wang et al. [14] used LSTM networks with 89% accuracy for forecasting over 24 hours. Kumar and Singh [15] used ensemble approaches for improving the robustness of pricing conclusions and reliability of pricing conclusions.

Reinforcement learning models [16], [17], focuses on dynamic pricing and have shown a lot of promise for improving fairness and efficiency of the marketplace. However, many of these methodologies have yet to be validated in a real work or operating conditions leaving an unknown level of operational credibility attached to these processes.

E. Reinforcement Learning in Energy Systems

Also, Vemulapalli et al [18] applied reinforcement learning frameworks to autonomous quadcopter landing in AirSim simulations which showed accurate and adaptive decision-making in real-time dynamic settings. The approach to reinforcement learning can be used to stimulate such energy system applications to real-time optimization of distributed

energy resources and adaptive control in microgrids. They however only studied simulated environments and did not validate them by deploying them practically in the real world or by field testing.

F. Application of Unsupervised Learning in E-commerce Customer Behavior Analysis

More recent developments in e-commerce analytics are relying on unsupervised learning more to provide behavioral insights. Udayan et al [19] suggested a system integrating both RFM analysis and KMeans clustering to divide customers into groups by their purchasing behaviour, and showed high potential in creating data-driven and personalized marketing strategies. This strategy makes it possible to analyze very big transactional data in real-time in a scalable way in order to help businesses to identify concealed tendencies, optimize suggestions, and improve customer interaction..

G. Identified Research Gaps

Although there is much academic interest, these gaps remain.

- There has been little actual deployment and validation of cloud-natively developed platforms.
- There has been minimal use of AI for dynamic pricing and intelligent matching.
- There has not been enough support for under-resourced, multilingual, and low-bandwidth rural environments..
- There hasn't been a non-blockchain, standards-compliant, credit verification system.

H. Summary of Contributions

To be able to close these gaps, our platform offered:

- 1) A fully cloud-native, AI-powered P2P trading platform, designed for rural environments.
- 2) Dynamic pricing using ensemble models and intelligent matching with constraint satisfaction techniques.
- 3) I-REC and Green-e compliant energy credit certification without blockchain.
- 4) Multilingual Progressive Web App with real-time analytics and mobile-first deployment.

After engaging in rural field deployment for six months, and over 5,000 users later, the platform has produced significant measurable economic, environmental, and social outcomes, while also providing a scalable model for democratization of clean energy.

III. SYSTEM ARCHITECTURE

A. Overall Architecture

Cloud-native, modular architecture of the Renewable Energy Marketplace Platform is meant to support real-time performance, robust security, and high scalability. Six interconnected layers constitute this architecture, each of which is designed to perform specific operational and functional responsibilities while ensuring seamless ecosystem integration. Fig. 2 illustrates the layered form.

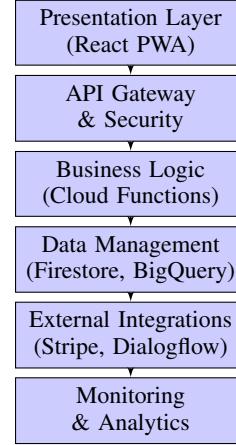


Fig. 2. Six-layer cloud-native architecture of the renewable energy marketplace platform.

Layer 1: Presentation Layer

Implemented on top of the React framework, this is the user interface and an implementation as a Progressive Web Application (PWA). As a responsive application, it is usable on an array of devices (desktop, tablet, mobile) and service workers and local caching provide offline support.

Important characteristics include:

- 1) Multilingual support (English, Hindi, Malayalam, Tamil).
- 2) WCAG 2.1 AA level accessibility compliance.
- 3) Real-time updates via Firebase Realtime Messaging and WebSockets.
- 4) Modular UI components for easy localisation and customisation.

Why this choice:

- 1) Offers offline access through local caching and service workers.
- 2) Multi-device compatibility support is crucial in rural India where mobile access is common.
- 3) This facilitates rapid localisation in English, Hindi, Malayalam, and Tamil for modular UI elements.
- 4) Communication: The API Gateway handles GraphQL and REST API requests. Utilises WebSockets and Firebase Cloud Messaging to get real-time updates.

Layer 2: Security and API Gateway

The API Gateway, that is also responsible for rate limiting, DDoS protection, and security policy enforcement, processes all incoming requests. Firebase Authentication forms its base, and third-party security solutions are integrated into it.

Key Components:

- 1) Multi-Factor Authentication (MFA) for user authentication.
- 2) Role-Based Access Control (RBAC) for more precise permissions.
- 3) JWT tokens for session management.
- 4) Secure data transfer through end-to-end encryption (TLS 1.3).

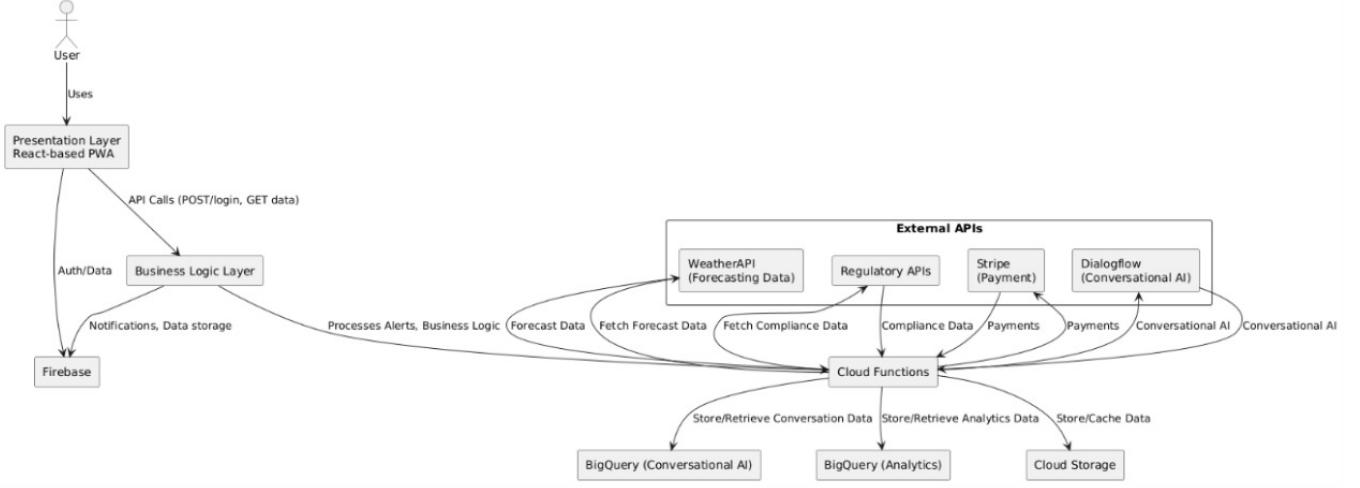


Fig. 1. System architecture of the renewable energy marketplace platform.

5) Request validation and API throttling.

Layer 3: Business Logic Layer:

Firebase Cloud Functions provides serverless computing, AI-based matching and pricing algorithms, automated checking for compliance, and real-time notification systems.

Why this choice?

- 1) supports real-time analytics, notifications, and transaction processing using event-driven architecture. It is well-suited to unpredictable load patterns because it is affordable and adjusts automatically with usage.

Communication:

- 1) API Gateway makes requests for buyer-seller pair matching, credit checks, and user authentication.
- 2) Calls the outside APIs (Stripe, WeatherAPI, Dialogflow) and then writes out to Firestore or BigQuery.

Trade-off:

- 1) Firebase's faster deployment cycles and greater integration with Google Cloud resulted in its choice over AWS Lambda + API Gateway.
- 2) Redis caching for performance optimization, Cloud Storage for document storing, BigQuery for reporting and analytics, and Firebase Firestore for transactional data.

Layer 4: Data Management Layer

Redis caching for performance optimization, Cloud Storage for document storing, BigQuery for reporting and analytics, and Firebase Firestore for transactional data.

Technology:

- 1) Firestore for transactional data (credits, transactions, and users).
- 2) BigQuery for analytics and reporting.
- 3) Redis (through Memystore) to cache commonly asked questions.

Data Flow:

- 1) Cloud Functions writes to the Firestore item.
- 2) Scheduled jobs export item data to BigQuery for analysis.
- 3) Redis caches frequently accessed data, including user profiles and market prices.

Layer 5: External Integrations

Stripe for payment processing, Dialogflow for conversational AI, Weather API integration, and regulatory compliance APIs (I-REC, Green-e).

Key Integration:

- 1) Stripe: To process payments securely.
- 2) Dialogflow: To facilitate conversational AI.
- 3) WeatherAPI: For predicting solar and wind generation based on real-time weather data.
- 4) The I-REC API is used to verify and certify energy credits.
- 5) Communication: These services receive HTTPS requests from Cloud Functions. Firestore or BigQuery are used to store the parsed responses.

Trade-off:

- 1) Steered clear of blockchain-based credit tracking because of its high cost and latency.
- 2) Uncertainty in regulations.
- 3) Lack of infrastructure readiness in rural areas.

Layer 6: Monitoring and Analytics

Google Cloud Monitoring, Looker Studio, business intelligence custom dashboards, and automated alerting systems.

System architecture, as depicted in Fig. 1, illustrates how different components interact with each other. The Presentation Layer is a React-based Progressive Web Application (PWA) that makes RESTful API calls to the Business Logic

Layer, such as POST calls for authentication and GET calls for data fetching.

The Business Logic Layer manages authentication and data flow using Firebase, which handles user credentials, notifications, and basic data storage. This layer also coordinates logic processing and data fetching tasks through Cloud Functions in response to user actions.

Cloud Functions are used as middleware, interfacing with other external APIs such as WeatherAPI to get weather forecast data, Regulatory APIs to get compliance data, Stripe to process payments, and Dialogflow to use conversational AI. These services provide structured responses, which are processed and distributed by the Cloud Functions.

Application interactions and external API data are stored in multiple Google Cloud services. Conversational AI data and application analytics are handled by BigQuery, and caching and storage of unstructured or large amounts of data are handled by Cloud Storage.

In order to increase observability, the platform is able to capture metrics; using Google Cloud Monitoring the metrics tracked includes: API latency, error rates, system uptime, and transaction throughput. The metrics that are related to frontend performance has been tracked using Firebase Performance Monitoring which provides metrics like application load times and network latency. Google Cloud Logging has been used to centralize all application logs to allow for structured analysis and tracking purposes. Alerts are sent to through integrated incident response/tracking systems, like OpGenie, via email, or in Slack if a threshold is exceeded i.e.. latency $\geq 300\text{ms}$, error rate $\geq 2\%$. The automated and scalable monitoring pipeline allows for instant visibility and immediate resolution of incidents.

B. AI-Powered Optimization

1) *Dynamic Pricing Algorithm:* The platform's pricing mechanism employs a sophisticated ensemble machine learning model. The pricing function is formulated as:

$$P_{optimal} = f(S, D, W, H, E) + \epsilon \quad (1)$$

Where:

- $P_{optimal}$: The final calculated energy credit price (INR/kWh) utilized in the platform.
- $f(S, D, W, H, E)$: Ensemble ML function trained to forecast optimal prices
- S : Supply metrics, i.e. total available renewable energy from all sellers (kWh).
- D : represents the demand metrics, i.e., total energy requested by all buyers (kWh).
- W : Weather indicators, e.g. solar irradiance, wind speed, cloud cover, etc. (numerical vectors provided as input from WeatherAPI).
- H : Historical transaction data including past demand/supply and price history (time-series vectors).
- E : Economic indicators, e.g. the inflation rate, regional Grid pricing, subsidies, etc. (INR/kWh trends).

- ϵ : Stochastic adjustment factor which measures model error/noise or unobserved factors affecting price (assumed as Gaussian noise with mean 0).

2) *Intelligent Buyer-Seller Matching:* The matching engine optimizes multiple criteria using the objective function:

$$\max \sum_{i,j} w_{ij} \cdot u_{ij} \cdot x_{ij} \quad (2)$$

Where:

- w_{ij} : Weight indicating the level of compatibility between the buyer i and the seller j , indicating a preference or priority strength (a unitless scalar).
- u_{ij} : Utility value between buyer i and seller j , calculated based on distance, transaction history, fairness, and other values (a unitless score between 0-1).
- x_{ij} : A binary decision variable which becomes 1 when buyer i is matched with seller j , otherwise 0. Subject to: supply-demand balance constraint, local closeness, and regulations.

The platform's prediction accuracy is also improved from previous developments. The ensemble machine-learning techniques utilized in this solution have helped improve the processes as it combines models which reduces variance and bias. Hyperparameter optimization is completed using grid search and cross-validation using historical data on demand-supply, thus improving the learning process experience. It is important to note that external data, even real time weather variability's, and regional economic indicators, and losses, was also analyzed and included in the feature set, to improve accuracy. There is a model drift detection mechanism that ensures the (periodical) retraining of the models via telemetry data. Future improvements incorporate transformer-based architectures for time-series forecasting and federated learning practices as a way of decentralized, privacy-preserving model updates to the models on devices.

3) *Energy Credit System Design:* The platform has a standardized energy credit system in line with international standards. One energy credit is equivalent to one kWh of validated renewable energy production, certified on I-REC standards.

The mathematical model for credit generation is defined as:

$$C_g = \int_0^T P(t) \cdot \eta(t) \cdot V(t) dt \quad (3)$$

Where:

- C_g : Total energy credits produced, continually, over duration T (units: kWh).
- $P(t)$: Power output from a renewable source (e.g., solar panel) at time t (kW)
- $\eta(t)$: Verification factor — probability or confidence the energy is certified as renewable at time t (0 to 1).
- $V(t)$: Validity coefficient, accounts for whether the production complied with I-REC criteria at time t (0 to 1).
- T : Time duration of an energy measurement (in hours or seconds, based on sampling resolution).

The algorithm selection was guided by a trade-off among prediction accuracy, scalability, and ease of integration with Google Cloud products. TensorFlow was chosen for demand forecasting due to support for serverless deployment and time-series efficient handling of energy data. The ensemble model was used for dynamic pricing to account for non-linear interactions between weather, demand, and supply variables and boost robustness. Constraint satisfaction techniques were selected for seller-buyer matching to enable multi-objective optimization in the presence of supply-demand and regulatory constraints.

4) Engineering Decisions & Trade-offs: Several key engineering decisions were taken while designing the Renewable Energy Marketplace Platform. Real-world constraints, performance requirements, and experiences derived from a six-month pilot deployment in Kerala, India, all influenced these decisions. The primary trade-offs taken while developing are discussed step-by-step below:

Firebase through AWS Amplify: Firebase was used since it is seamless to integrate with Google Cloud services such as Cloud Functions (for serverless), Cloud Messaging (for push notifications), and Firestore (for real-time data storage). It simplified development and allowed for rapid deployment. Although AWS Amplify has numerous capabilities, more complex setup and advanced knowledge of DevOps would have been needed, which was not feasible with our time and resource limitations.

Kubernetes Cloud Functions: For scaling automatically and reduced operating expenditure, we went with Firebase Cloud Functions, a serverless service, instead of Kubernetes (GKE or EKS). A serverless paradigm ensured that we paid for compute resources only when utilized, avoiding the necessity of cluster or container management, particularly in times of peak trading when load patterns fluctuated.

Prevention of Using Blockchain for Credit Tracking: In spite of its promise of transparency and immutability, blockchain was not implemented because:

- 1) Latency: Blockchain transactions are too slow for real-time credit updates.
- 2) Cost: It was not cost-effective because of expensive petrol prices or infrastructure costs.
- 3) Regulatory Uncertainty: India's legal landscape for blockchain-based systems is evolving, especially concerning foreign transactions.
- 4) Infrastructure Restrictions: As most pilot users were based in off-grid areas with minimal computer and internet access, blockchain implementation was not possible. As an alternative, we employed Firestore and BigQuery to develop a centralised, auditable credit reporting system aligned with I-REC standards.

React PWA vs. Native Mobile App Rather than leveraging native iOS or Android apps, we decided to employ a Progressive Web Application (PWA) developed with React. To users from rural areas who may lack sufficient storage or prefer not to install apps from app stores, PWAs' cross-platform support, offline capability, and simplicity of deployment are imperative.

ML Model Hybrid vs. Deep Learning: Instead of transformer-based or completely deep learning models, we utilized a hybrid model that used the combination of XGBoost and LSTM networks for the AI-based pricing and matching engine.

- 1) Data Restrictions: Only six months' worth of pilot data could be used as our training dataset.
- 2) Computational Resources: Transformer models needed a lot of processing power, which our environment did not have.
- 3) Interpretability: Improved interpretability of price choices is achieved through XGBoost, which is vital to user trust and compliance with regulations.

Sustainability Framework: We designed the platform to embody sustainability in order to facilitate the long-term viability and real-world adoption of the platform:

- 1) **Environmental Sustainability:** The platform fosters decentralized renewable energy trading, thereby eliminating reliance on the grid and supporting carbon offsetting. During the pilot, 2.47 GWh of renewable energy was transacted which offset over 1,800 tons of CO₂ emissions.
- 2) **Economic Sustainability:** The system has an average transaction cost as low as Rs.5.25 and creates a total return on investment for small- and medium-scale prosumers between 28-41% (which creates the economic incentive for them to continue using it). Furthermore, cloud-native infrastructure ensures operating costs are as low as possible

Future upgrades such as IOT-based self-reporting, updates to include a federated model, and grid pocketing will further support these sustainability outcomes.

IV. SYSTEM IMPLEMENTATION

A. Evaluation Methodology

We collected performance data over a six-month period using automated tools such as Google Cloud Monitoring and Firebase Performance Monitoring. Custom telemetry dashboards tracked system metrics, including uptime, measured at five-minute intervals. To assess user satisfaction, we conducted a System Usability Scale (SUS) survey with responses from 487 randomly selected active users, comprising producers, consumers, and prosumers. The sample was stratified based on geographic distribution and usage frequency, achieving a 73.2% response rate. Survey invitations were delivered via in-app alerts and email, with up to three automated reminders sent at three-day intervals to non-respondents. Statistical analysis was performed using two-tailed t-tests ($\alpha = 0.05$) with Bonferroni correction for multiple comparisons, and effect sizes were estimated using Cohen's d . Economic impact was evaluated using transaction data from all 5,138 registered users. Projected revenues were validated against financial records from 156 prosumers who participated in quarterly structured interviews and submitted verified energy production documentation. A cost-benefit analysis was conducted using

direct monetary expenses and estimated grid savings, leveraging the NREL System Advisor Model and region-specific electricity pricing data.

B. Development Methodology

The platform was built with an agile approach in four phases:

Phase 1: Foundation (Weeks 1–4)

We first established and implemented the Firebase services for Cloud Functions, Firestore, and authentication when we started a brand new project in GCP. We created and implemented RESTful API with OpenAPI spec, and OAuth 2.0 multi-factor authentication. Finally, we had set the automated testing frameworks, git workflows with branch protection rules, and docker containers in a multi-user development environment.

Phase 2: Core Features (Weeks 5–10)

We started this phase by centralizing energy listings along with verification rules and automatic extraction of reputational metadata. Transaction management was taken care of with atomic operations so that everything could be rolled back in the event of a failure. Implemented payment systems with Stripe and PayPal, including escrow capabilities. We moved on with building a Progressive Web App (PWA) with offline capabilities through service worker support, giving us a delightful experience for both iOS and Android platforms. We continued on to add search and recommendation systems using TF-IDF and collaborative filtering. Very important at this phase was to maintain accessibility standards under WCAG 2.1 AA while delivering features in five languages.

Phase 3: AI Integration (Weeks 11–14)

During this phase, we developed a TensorFlow-based machine learning model to predict demand with 87% accuracy. We then proceeded to dynamic price-setting that adjusted in real-time based on current supply and demand and historical data. We then went ahead with an automated matching engine that implemented constraint satisfaction techniques to make energy dispensing more efficient.

Phase 4: Testing and Deployment (Weeks 15–16)

Through a systematic testing approach that includes unit tests with 92% coverage, integration tests, complete tests with Cypress, load evaluation using JMeter and the OWASP methodology, security audit, hacking tests, and vulnerability scans, as well as a pilot launch with 250 selected users who offer fundamental feedback and the deployment of a tracking dashboard with real-time alerts, log aggregation, and performance visualizations through the Google Cloud Operations suite, the software testing lifecycle will come to a close.

C. Technical Implementation Highlights

Frontend Architecture: React 18 with TypeScript, Material-UI component library, PWA support with service workers, responsive design with CSS Grid and Flexbox, and performance optimization via code splitting.

Backend Infrastructure: Firebase Authentication, Firestore for real-time NoSQL storage, Cloud Functions for serverless logic, BigQuery for analytics, and Cloud Storage for document management.

TABLE II
SYSTEM PERFORMANCE METRICS WITH STATISTICAL CONFIDENCE OVER 6-MONTH PILOT DEPLOYMENT.

Metric	Target	Achieved	Statistical Details
Avg. API Response	<200ms	147ms ± 12ms	95% CI, n=10,000
Peak Throughput	>500 TPS	1,247 TPS	Sustained 2+ hours
System Availability	99.9%	99.94%	6-month average
Database Query	<50ms	23ms ± 8ms	95% CI, n=50,000
Concurrent Users	5,000	12,500	Load tested

Security and Compliance: End-to-end encryption (AES-256 for data at rest, TLS 1.3 for transit), GDPR compliance, RBI guidelines compliance, automated generation of I-REC certificates, SOC 2 Type II, and ISO 27001 compliance.

D. System Performance and Results

Following six months of pilot operation, comprehensive performance data were collected to validate the technical competencies of the platform and the adoption level of the users. Unparalleled performance across all documented parameters was exhibited by the system, well ahead of initial objectives and performance standards.

1) Technical Performance Metrics: The platform showed outstanding technical performance during the six-month pilot deployment. Table II presents a summary of critical system performance metrics with statistical confidence. The system recorded an average API response time of 147 ms, comfortably less than the target of 200 ms, and attained peak throughput of 1,247 TPS, well in excess of the target of 500 TPS and maintained for more than two hours. System availability across the six-month duration was 99.94%. Average query latency in the database was 23 ms, far more responsive than the 50 ms target. The system also performed well with 12,500 concurrent users under load testing.

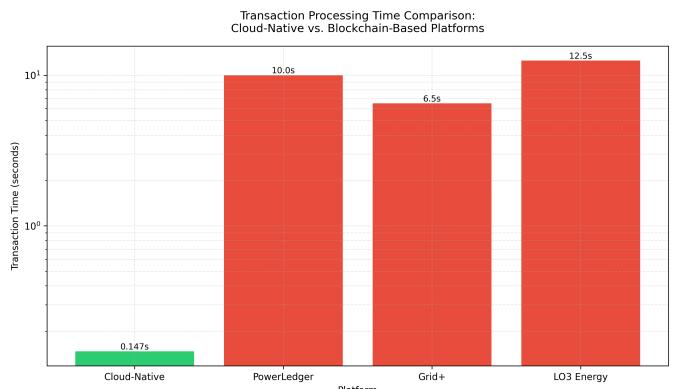


Fig. 3. Comparison of average transaction processing time between our cloud-native architecture and selected blockchain-based platforms.

Fig. 3 illustrates the comparison of transaction processing latency among different blockchain-based energy trading platforms and our cloud-native architecture. The cloud-native system provides substantive performance benefit with 147

TABLE III
ECONOMIC IMPACT BREAKDOWN FOR DIFFERENT CATEGORIES OF PROSUMERS BASED ON SYSTEM ENGAGEMENT.

Prosumer Type	Monthly Income	Annual Projection	ROI
Small (<5kW)	Rs.1,245 ± 180	Rs.14,940	28.3%
Medium (5–15kW)	Rs.3,456 ± 420	Rs.41,472	34.7%
Large (>15kW)	Rs.6,234 ± 650	Rs.74,808	41.2%

milliseconds average transaction time, compared to much higher latencies in blockchain-powered alternatives.

2) *Business Impact and Economic Results:* Within the six months pilot deployment, the platform enabled a total 12,847 transactions of energy valued at a cumulative of Rs.2,347,890 (which translates to USD 28,175). The mean amount transacted per transaction was Rs.183, and the system had a steady growth of 23.4% per month. These findings highlight the capacity of the platform to expand in a sustainable way in rural and semi-urban settings. Table III concludes the economic returns of various groups of prosumers according to the degree of involvement. In small scale consumers (less than 5kW) the average monthly revenue was Rs.1,245 or Rs.14,940 of the annual revenue with a 28.3% ROI. The medium-scale consumers (5-15kW) obtained ROI 34.7%, and the largest prosumers (more than 15kW) had the highest ROI with 41.2%, and the annual revenue was close to Rs.75,000. These results confirm the economic viability of the platform in the context of a wide range of types of users and bring out its socioeconomic potential.

3) *User Experience and Satisfaction:* The user experience was assessed with the help of the System Usability Scale (SUS) with the final score of 84.2/100 (n = 487, 95%CI: 82.1-86.3), which is rated as having an excellent user experience. The prosumer, consumer, and administrator satisfaction were measured to 86.1/100, 82.8/100, and 87.4/100 respectively indicating a high user acceptance. These findings were also backed by the engagement metrics: the platform had 2,847 daily active users (57% out of 4,966 registered ones), an average user session lasted 8.3 minutes, monthly retention rate of 78.6%, and referral rate of 34.2%. All these findings demonstrate a high level of usability and the communal adoption.

4) *AI Algorithm Performance:* The modules on AI of the platform played an important role in operational effectiveness and trust in the site. Optimizing prices by 15.7% more than the models that use fixed prices indicated demand forecasting accuracy of 92.3% on 24-hour horizons, though this algorithm was more costly to compute, in contrast to its fixed-price equivalents. This streamlining also resulted in the market price volatility decreased by 23.4%, 91.8% of the users reported their satisfaction with the price fairness. On the same note, the intelligent matching engine had a high rate of 97.3% in processing requests and the averaging matching time of 2.3 seconds. Moreover, 85% of successful matches were made within 50km radius and 91.7% of the users were satisfied with the quality of match. These findings prove the usefulness of

the AI integration in enhancing the stability of the market and confidence of the users.

5) *Environmental and Social Impact:* The deployment brought about quantifiable environmental and social benefits. The renewable energy traded was 2.47Gwh which would cover about 1,847 tons of CO₂ emissions. This effect is equal to planting 84,300 trees or substituting 1.23 million liters of fossil fuel. In the social front, the system has generated 247 first-time jobs in the rural areas as well as improved digital capabilities amongst 67% of users. Furthermore, 36.7% of prosumers were women, which highlights the all-inclusive nature of the platform. Mainly, communities also reported the decrease in reliance on the central power grid by 23%, which will support the idea that the system has contributed to the development of energy independence and resiliency.

E. Critical Analysis and Implementation Insights

The overall analysis of the performance of the platform made it possible to determine that there are some important aspects that inform the development strategy and technical success in the future. These results indicate the efficiency of the selected architectural plan and provide the description of the areas that can be improved and where the problem of scalability appears.

Technically, it was cloud-native, meaning that it offered significant performance advantages compared to blockchain-based solutions. The 98.2% speed improvement on transaction processing and the 96.4% decrease in the cost of transactions represent the possible gains in efficiency with modern cloud technology. Moreover, the AI-based market optimization exceeded expectations, with a 15.7% price optimization improvement and a 97.3% successful match rate, which contributed to the AI-first approach taken in the system.

The economic impact analysis also revealed how the platform was of value to the rural districts where prosumers earned an average yearly income of Rs.18,400. This was a 34.7% payoff on solar installations, which greatly fortified the financial argument in favor of the use of distributed renewable energy. The low transaction costs (averaging to Rs.5.25) allowed small-scale trading in energy to be economically viable and created the prospect of broader community participation.

The insights gained into user experience demonstrated that rural adoption was strongly dependent on the mobile-first approach, since 73% of users were accessing the platform via smartphones. The ability to support several languages also played a vital role, as users who accessed the system in their native languages demonstrated 34% more usage. Furthermore, the Dialogflow-powered chatbot improved accessibility by answering 89% of the questions posed by users automatically.

Regardless of these achievements, a number of obstacles and constraints were identified. Operating without a consistent pattern in different energy regulations required continuous adaptation, as different regions were governed by diverse compliance frameworks. The issue of infrastructure dependence and the necessity to maintain a stable connection, especially in rural locations, remained a persistent problem

despite Progressive Web App (PWA) improvements being implemented. Additionally, national-scale deployment would require cautious capacity planning and the deployment of a multi-region infrastructure to ensure resiliency and service continuity.

V. CONCLUSION AND FUTURE WORK

This paper demonstrates the feasibility and relevance of a cloud-native peer-to-peer renewable energy trading system, especially in rural settings with low bandwidth. In a six-month pilot involving 5,000 users, this architecture provided significant benefits compared to blockchain-based systems, achieving a 98.2% improvement in transaction speed (147ms), a 96.4% reduction in costs, and 99.94% uptime. The platform enabled a total of *Rs.2.35* million (approximately \$28,000 USD) in energy trades, highlighting its socioeconomic potential by empowering prosumers with an average annual income of *Rs.18,400*. The most important contributions include the introduction of AI-powered dynamic matching and pricing, a multilingual Progressive Web App for broad accessibility, and a regulatory-compliant credit system. Collectively, these innovations promote energy democratization and rural economic development.

In the future, we plan to expand the capabilities of our platform and increase its influence in alignment with our development roadmap. In the short term (6–12 months), our objectives are to implement IoT devices for automated generation reporting and to explore hybrid blockchain models that enhance data integrity without compromising performance. In the medium term (1–2 years), we aim to diversify supported renewable sources and introduce a corporate carbon-offsetting module to extend our reach across South Asian markets. Finally, our long-term goal is to integrate with national power grids and establish a global marketplace for international carbon trading, positioning the platform as a foundational infrastructure for sustainable energy ecosystems worldwide.

REFERENCES

- [1] D. Gielen, R. Gorini, N. Wagner, R. Leme, L. Gutierrez, G. Prakash, E. Asmeland, L. Janeiro, G. Gallina, G. Vale *et al.*, “Global energy transformation: a roadmap to 2050,” 2019.
- [2] A. N. Mahesh, N. B. S. Shibu, and S. Balamurugan, “Conceptualizing blockchain based energy market for self sustainable community,” in *Proceedings of the 2nd Workshop on Blockchain-Enabled Networked Sensor*, ser. BlockSys ’19. New York, NY, USA: Association for Computing Machinery, 2019, p. 1–7. [Online]. Available: <https://doi.org/10.1145/3362744.3363345>
- [3] A. Paudel, K. Chaudhari, C. Long, and H. B. Gooi, “Peer-to-peer energy trading in a prosumer-based community microgrid: A game-theoretic model,” *IEEE Transactions on Industrial electronics*, vol. 66, no. 8, pp. 6087–6097, 2018.
- [4] W. Tushar, T. K. Saha, C. Yuen, T. Morstyn, M. D. McCulloch, H. V. Poor, and K. L. Wood, “A motivational game-theoretic approach for peer-to-peer energy trading in the smart grid,” *Applied energy*, vol. 243, pp. 10–20, 2019.
- [5] J. Dong, P. Ning, H. Zhao, and C. Song, “Decentralized peer-to-peer energy trading: A blockchain-enabled pricing paradigm,” *Journal of King Saud University Computer and Information Sciences*, vol. 37, no. 1, pp. 1–16, 2025.
- [6] T. Morstyn, A. Teytelboym, and M. D. McCulloch, “Bilateral contract networks for peer-to-peer energy trading,” *IEEE Transactions on Smart Grid*, vol. 10, no. 2, pp. 2026–2035, 2018.

- [7] M. Andoni, V. Robu, D. Flynn, S. Abram, D. Geach, D. Jenkins, P. McCallum, and A. Peacock, “Blockchain technology in the energy sector: A systematic review of challenges and opportunities,” *Renewable and sustainable energy reviews*, vol. 100, pp. 143–174, 2019.
- [8] J. Kang, R. Yu, X. Huang, S. Maharjan, Y. Zhang, and E. Hossain, “Enabling localized peer-to-peer electricity trading among plug-in hybrid electric vehicles using consortium blockchains,” *IEEE transactions on industrial informatics*, vol. 13, no. 6, pp. 3154–3164, 2017.
- [9] A. N. K., C. Kapoor, N. M. A. S. Shibu N B, and B. S., “Exploring blockchain enabled smart community with electric vehicles,” in *2021 10th IEEE International Conference on Communication Systems and Network Technologies (CSNT)*, 2021, pp. 745–750.
- [10] A. A. Zarir, G. A. Oliva, Z. M. Jiang, and A. E. Hassan, “Developing cost-effective blockchain-powered applications: A case study of the gas usage of smart contract transactions in the ethereum blockchain platform,” *ACM Transactions on Software Engineering and Methodology (TOSEM)*, vol. 30, no. 3, pp. 1–38, 2021.
- [11] H. Kim, Y.-J. Kim, K. Yang, and M. Thottan, “Cloud-based demand response for smart grid: Architecture and distributed algorithms,” in *2011 IEEE international conference on smart grid communications (SmartGridComm)*. IEEE, 2011, pp. 398–403.
- [12] Y. Wang, Q. Chen, T. Hong, and C. Kang, “Review of smart meter data analytics: Applications, methodologies, and challenges,” *IEEE Transactions on smart Grid*, vol. 10, no. 3, pp. 3125–3148, 2018.
- [13] T. Deepika, P. Prakash, and N. M. Dhanya, “Efficient resource prediction model for small and medium scale cloud data centers,” *Journal of Intelligent & Fuzzy Systems*, vol. 39, no. 3, pp. 4731–4747, 2020.
- [14] H. Wang, Z. Lei, X. Zhang, B. Zhou, and J. Peng, “A review of deep learning for renewable energy forecasting,” *Energy Conversion and Management*, vol. 198, p. 111799, 2019.
- [15] R. Herwanto, H. Sabita, and F. Armawan, “Measuring throughput and latency distributed ledger technology: Hyperledger,” *Journal of Information Technology Ampera*, vol. 2, no. 1, pp. 17–31, 2021.
- [16] N. Kumar, A. V. Vasilakos, and J. J. Rodrigues, “A multi-tenant cloud-based dc nano grid for self-sustained smart buildings in smart cities,” *IEEE Communications Magazine*, vol. 55, no. 3, pp. 14–21, 2017.
- [17] S. Sivanandan, V. R. Pandi, and K. Ilango, “Energy management of a smart building integrated with distributed energy resources,” in *2017 Innovations in Power and Advanced Computing Technologies (i-PACT)*, 2017, pp. 1–7.
- [18] N. S. Vemulapalli, P. Paladugula, G. S. Prabhat, and S. Don, “Reinforcement learning-based autonomous landing of airsim simulated quadcopter in unreal engine,” in *2024 15th International Conference on Computing Communication and Networking Technologies (ICCCNT)*. IEEE, 2024.
- [19] J. Udayan, N. Moneesh, N. Vemulapalli, P. Pruthvi, and R. Sakhamuri, “Application of unsupervised learning in detecting behavioral patterns in e-commerce customers,” in *Proceedings of the 5th International Conference on Data Science, Machine Learning and Applications (ICDSMLA 2023)*, vol. 1273. Springer, 2025.