

# Secure Reinforcement Learning-Based Time Synchronization: Defending Against Adversarial Attacks with Anomaly Detection

Karthika Sreenivasan, Malavika Ajith, Krishnpriya J D

Department of Computer Science and Engineering, Amrita Vishwa Vidyapeetham, Amritapuri, India

Roll No.: {AM.EN.U4CSE22237, AM.EN.U4CSE22242, AM.EN.U4CSE22235}

**Abstract—Purpose:** Accurate time synchronization is vital for consistency and coordination in distributed systems. However, adversarial attacks—such as NTP spoofing, Sybil, and Byzantine faults—pose serious threats by manipulating timestamps, impacting systems like blockchain and cloud services. To address this, we propose the Secure Reinforcement Learning-Based Adaptive Time Synchronization (SRL-TS) framework. It combines a Reinforcement Learning agent with an autoencoder-based anomaly detector to dynamically adjust sync intervals and detect malicious time updates. A Trust Score System evaluates node reliability to filter out attackers, enhancing resilience and accuracy.

**Evaluation:** SRL-TS was tested under simulated attack scenarios and demonstrated improved synchronization accuracy and robustness. It successfully detects malicious nodes, adapts sync intervals, and maintains system reliability, making it an effective and scalable solution for secure time synchronization in distributed networks.

**Index Terms**—Distributed Systems, Time Synchronization, Reinforcement Learning, Anomaly Detection, Trust Score, Adversarial Attacks

## I. INTRODUCTION

In distributed systems, accurate time synchronization is a foundational requirement. Time-stamped operations are central to many critical functions including data consistency, event ordering, coordination, resource sharing, and system recovery. Any deviation in synchronized time across nodes can result in inconsistencies, data loss, or security vulnerabilities. As these systems expand across diverse and dynamic network environments, maintaining consistent and secure time synchronization becomes increasingly complex.

Traditional time synchronization protocols such as the Network Time Protocol (NTP) are vulnerable to a range of Time Synchronization Attacks (TSAs), including spoofing, Sybil, and Byzantine faults. These attacks manipulate the timing information exchanged between nodes, disrupting the logical sequence of operations and potentially destabilizing systems like financial services, blockchain infrastructure, and cloud-based platforms. Furthermore, most existing systems lack adaptability to network changes and fail to detect or mitigate adversarial behavior in real-time.

To address these challenges, we propose the Secure Reinforcement Learning-Based Adaptive Time Synchronization (SRL-TS) framework. This system leverages Reinforcement Learning (RL) to dynamically adjust synchronization intervals based on current network conditions and uses anomaly

detection to identify suspicious or inconsistent timestamps. By introducing a Trust Score System, it also penalizes malicious nodes and ensures that only reliable sources influence the synchronization process.

This report presents the design, implementation, and evaluation of the SRL-TS framework. The subsequent sections cover a review of related work, the specific problem definition, the architecture and algorithms used, implementation details, experimental results, and the overall conclusion and future scope of the project.

## II. RELATED WORKS

Time synchronization plays a vital role in ensuring correctness and consistency across distributed systems. Classic systems such as the Google File System (GFS) and the Hadoop Distributed File System (HDFS) rely heavily on accurate timestamps for managing distributed storage, replication, and failure recovery. These systems often use basic synchronization mechanisms like the Network Time Protocol (NTP), which are vulnerable to spoofing attacks if not properly secured.

Foundational works like Lamport's logical clocks and the Chandy-Lamport snapshot algorithm provide logical ordering of events and consistent global state capturing in distributed environments. However, they operate under the assumption of trustworthy nodes and do not handle adversarial behaviors such as Sybil or Byzantine attacks.

Zhang et al. investigated vulnerabilities in NTP and demonstrated how timestamp spoofing can destabilize distributed systems. While their solution outlined security patches, it lacked adaptability and dynamic response to changing network conditions.

Bahl et al. proposed a blockchain-based timestamping mechanism to enhance data integrity in distributed systems. Although secure, this approach incurs high latency and computational overhead, making it unsuitable for lightweight or real-time applications.

Wang et al. explored the use of Reinforcement Learning (RL) to optimize synchronization intervals in wireless sensor networks. Their model, based on the PPO algorithm, achieved adaptive timing but lacked any mechanism to detect or respond to malicious behavior, limiting its robustness in untrusted environments.

Li and Chen developed an anomaly detection system using autoencoders for identifying irregular network patterns. While effective in flagging anomalies, their method was not integrated into any adaptive control framework and offered no countermeasures against detected threats.

Gong et al. introduced a static trust scoring system for peer-to-peer networks to isolate unreliable nodes. Although conceptually strong, their method did not dynamically adjust to changing node behavior or support real-time decision-making.

Chaudhuri et al. examined hybrid synchronization techniques that blend physical and logical clocks to improve synchronization accuracy. Their work demonstrated high precision in controlled environments but assumed cooperative node behavior and lacked built-in security measures.

While many individual works address synchronization accuracy, anomaly detection, or trust evaluation independently, very few integrate all three into a unified, adaptive system. RL-based solutions often ignore security; anomaly detection models typically do not influence system control; and trust systems are often static and fail to adjust in real-time.

The proposed Secure RL-Based Adaptive Time Synchronization (SRL-TS) framework addresses this critical gap by combining a PPO-based RL agent for adaptive interval control, autoencoder-driven anomaly detection for identifying timestamp manipulation, and a Trust Score System that evolves based on node behavior. This integration ensures the system not only adapts intelligently to network dynamics but also remains resilient under active threats such as spoofing, Sybil, and Byzantine attacks.

### III. PROBLEM DEFINITION

Accurate and consistent time synchronization is fundamental in distributed systems for coordinating processes, ensuring data consistency, and enabling effective fault detection. However, achieving reliable synchronization is especially challenging in adversarial environments where some nodes may act maliciously by sending false or manipulated timestamps. These malicious behaviors can significantly disrupt synchronization protocols, causing errors that propagate through the system and potentially leading to security vulnerabilities or system failures.

For instance, in real-world applications such as sensor networks for industrial automation or smart grids, incorrect timing due to malicious nodes can result in improper event ordering or trigger false alarms. This affects the overall safety and efficiency of the system. Conventional synchronization protocols typically assume nodes behave honestly and often operate on fixed synchronization intervals, leaving the system vulnerable to attacks such as spoofing, Sybil attacks, or Byzantine faults that compromise time accuracy.

The core problem to be solved is to develop a secure and adaptive time synchronization mechanism capable of dynamically adjusting synchronization intervals based on current network behavior and node trustworthiness. The system must detect and mitigate the effects of malicious nodes, ensuring

that synchronization remains accurate and reliable even in the presence of attacks and network uncertainties.

#### A. Constraints and Requirements

The solution must be scalable, efficiently supporting synchronization across multiple nodes without introducing excessive computational or communication overhead. It should be fault-tolerant, maintaining synchronization accuracy despite the presence of faulty or malicious nodes. Minimizing communication overhead is critical to avoid network congestion, meaning synchronization and anomaly detection processes must be lightweight and efficient.

The system also needs to adapt synchronization intervals in real time, responding to varying network conditions and anomaly detection results. Robustness against a range of attack types is essential, requiring mechanisms to detect, isolate, and ignore unreliable nodes. Finally, the algorithms must be resource-efficient to accommodate nodes with limited processing power and memory.

#### B. System Design Goals

The design aims to ensure robust and accurate time synchronization across all nodes, even when some behave maliciously. It seeks to incorporate adaptive mechanisms that intelligently adjust synchronization intervals for improved efficiency and resilience. Detecting and mitigating malicious activity through anomaly detection and trust scoring is a fundamental goal to preserve system integrity.

Another important goal is to minimize communication overhead, optimizing the synchronization protocol to reduce unnecessary network traffic. Lastly, the system should be scalable to large distributed networks and extensible, allowing future integration of additional attack models or synchronization enhancements without major redesigns.

### IV. METHODOLOGY

The Secure RL-Based Adaptive Time Synchronization (SRL-TS) framework is designed to provide robust and intelligent time synchronization in distributed systems under adversarial conditions. The methodology integrates Reinforcement Learning (RL), anomaly detection, and trust evaluation into a unified architecture that adapts to network dynamics while actively defending against attacks.

#### A. System Architecture

The system consists of multiple distributed nodes and a central controller. Each node maintains its local clock and periodically sends timestamp data to the controller. Nodes can either behave normally or act maliciously by sending manipulated timestamps to disrupt synchronization. The central controller processes incoming timestamps, evaluates node behavior using trust scores, detects anomalies through learned models, and makes synchronization adjustments accordingly. Key components include the node simulation environment, the controller module managing trust and anomalies, a reinforcement learning agent, and a trust scoring mechanism to filter unreliable nodes.

## B. Algorithms and Techniques

The core of the adaptive mechanism is a Reinforcement Learning agent trained with the Proximal Policy Optimization (PPO) algorithm. The agent observes network conditions and node trustworthiness to dynamically adjust synchronization intervals for optimal performance. Anomaly detection is performed by an autoencoder neural network trained to reconstruct normal timestamp patterns; deviations in reconstruction error signal suspicious activity. Trust scores for each node are continuously updated based on consistency and anomaly detection outcomes, with nodes falling below a trust threshold excluded from influencing synchronization decisions.

## C. Tools and Technologies

The SRL-TS framework is implemented in Python, leveraging OpenAI Gym to create the RL environment and train the PPO agent. Autoencoder models are built using TensorFlow or PyTorch. Socket programming facilitates multi-node timestamp communication over TCP, simulating distributed network behavior. Data manipulation and analysis are supported by NumPy, while Matplotlib and Seaborn are used for visualizing synchronization accuracy and trust score trends. This modular design promotes scalability and ease of extension for additional attack models or synchronization techniques.

## V. IMPLEMENTATION

### A. Step-by-Step Implementation

The implementation of the SRL-TS system began with simulating multiple distributed nodes, each maintaining a logical clock that periodically increments. These nodes communicate their local timestamps to a central controller using TCP socket connections. Some nodes were programmed to behave maliciously by sending manipulated timestamps to emulate adversarial attacks. The controller was implemented using non-blocking socket programming to handle simultaneous communications from multiple nodes efficiently. It collects and processes timestamp data, forwarding it to the anomaly detection engine and trust scoring system. The anomaly detection component, based on an autoencoder neural network, was trained on normal timestamp data to detect deviations indicative of malicious activity. Trust scores for each node were updated dynamically based on anomaly detection outputs and consistency checks, enabling the system to isolate untrustworthy nodes. A reinforcement learning agent using the Proximal Policy Optimization (PPO) algorithm was integrated within a custom OpenAI Gym environment to adaptively select synchronization intervals. The agent observes the current network state, including trust scores and timestamp deviations, and outputs the optimal synchronization timing to balance accuracy and communication overhead. Logging and visualization modules were added to monitor clock offsets and trust score trends during execution.

### B. Code Overview

Key parts of the implementation include socket-based communication protocols, the reinforcement learning training loop,

and the autoencoder model for anomaly detection. The socket programming ensures reliable message exchange between nodes and the controller, using asynchronous I/O to manage concurrency. The PPO agent's training loop involves interaction with the Gym environment, collecting rewards based on synchronization performance, and updating policy networks accordingly. The autoencoder is implemented with a simple feedforward architecture trained to reconstruct normal timestamp sequences; its reconstruction error serves as the anomaly metric. All code components are modularized with clear comments to enhance readability and maintainability.

## C. System Components

The system comprises several tightly integrated components. The client-server communication module uses TCP sockets to establish persistent connections and exchange timestamp messages. Fault handling is achieved through anomaly detection and trust scoring, which flag and isolate suspicious nodes to maintain synchronization integrity. The reinforcement learning module dynamically adjusts synchronization intervals in response to network conditions and detected anomalies, improving robustness and efficiency. Consistency across nodes is maintained by synchronizing trusted nodes' clocks based on the RL agent's recommendations. Data storage and visualization components enable analysis of system behavior and performance metrics such as offset accuracy and trust score evolution.

## D. Challenges Faced

During implementation, several challenges arose. Network latency and packet loss caused irregularities in timestamp reception, complicating anomaly detection and trust evaluation; these issues were mitigated through buffering strategies and timeout handling. Training the PPO agent required extensive hyperparameter tuning to ensure stable learning and prevent oscillations in synchronization intervals. Selecting an effective anomaly detection threshold was difficult, as too sensitive a threshold led to false positives while too lenient missed attacks; this was addressed through iterative cross-validation on representative datasets. Resource constraints on simulation hardware necessitated optimizing model sizes and processing batches to balance accuracy and efficiency. Despite these difficulties, the modular architecture facilitated debugging and incremental improvements.

## VI. RESULTS AND DISCUSSION

### A. Experiments

To evaluate the performance and effectiveness of the Secure RL-Based Adaptive Time Synchronization (SRL-TS) system, a series of experiments were conducted under both benign and adversarial conditions. The system was deployed in a simulated distributed environment consisting of multiple nodes, where a portion of the nodes were designated as malicious actors injecting false timestamps. The experiments aimed to test the adaptability, accuracy, and resilience of the time synchronization mechanism using several metrics including

clock offset deviation, synchronization frequency, trust score accuracy, and anomaly detection performance. The testing environment also introduced network noise and latency to emulate real-world uncertainties. Comparative baselines included fixed-interval synchronization and majority-vote-based consensus synchronization, against which SRL-TS was evaluated for improvement in trust-based filtering and adaptive interval control.

### B. Results

The results demonstrated that SRL-TS outperformed traditional synchronization methods in both accuracy and fault tolerance. In normal conditions, the average clock offset across all nodes was consistently maintained below a threshold of 2 milliseconds, while in adversarial settings, the system effectively isolated malicious nodes within the first 10 synchronization cycles based on trust score degradation. The anomaly detection engine achieved a detection accuracy of over 92

### C. Analysis

The experimental outcomes validate the design goals of SRL-TS, confirming its ability to maintain accurate time synchronization while defending against adversarial attacks. The dynamic adaptation of synchronization intervals via reinforcement learning reduced unnecessary communication, which is particularly beneficial for bandwidth-constrained environments. Compared to related approaches discussed earlier—such as fixed-interval or consensus-based synchronization—SRL-TS introduces a proactive and intelligent mechanism capable of adapting to changing network conditions and malicious behaviors. However, some limitations were observed. The autoencoder's performance declined slightly when trained on insufficient or biased data, indicating the importance of representative training datasets. Additionally, the RL agent required a non-trivial amount of training time and hyperparameter tuning, which may be a barrier to deployment in resource-constrained systems. Future improvements could include integrating lightweight models for edge environments, incorporating federated learning to preserve node privacy during RL training, and enhancing the explainability of trust score updates. Overall, the system proved effective and resilient, achieving synchronization with a high degree of reliability even under adverse network scenarios.

## VII. CONCLUSION

In summary, the Secure RL-Based Adaptive Time Synchronization (SRL-TS) system effectively addresses the challenge of maintaining accurate and robust time synchronization in distributed systems, even under adversarial conditions. The system's integration of reinforcement learning for dynamic synchronization interval adjustment, an autoencoder-based anomaly detection engine, and a trust scoring mechanism to isolate malicious nodes has resulted in significant improvements over traditional fixed-interval and consensus-based synchronization approaches. Experimental results indicate that

SRL-TS can maintain clock offsets within narrow bounds, quickly identify and mitigate the effects of manipulated timestamps, and optimize communication overhead. These findings underscore the importance of adopting intelligent and adaptive mechanisms in scenarios where security and precision are critical, such as sensor networks and smart grids.

Looking forward, several avenues for future work are apparent. Enhancing the scalability of the system to manage larger and more complex networks remains a key focus, potentially through distributed reinforcement learning approaches or hierarchical synchronization protocols. Future research could also explore alternative or hybrid machine learning models to further refine anomaly detection and trust scoring, thereby improving detection accuracy under diverse network conditions. Additionally, integrating federated learning techniques may help preserve privacy while distributing the training process across multiple nodes. Finally, further work on reducing the computational overhead and training time of the RL agent would facilitate deployment in resource-constrained environments, making the solution more adaptable and practical for real-world applications.

## REFERENCES

- [1] Tanenbaum, A. S., & van Steen, M. (2017). *Distributed Systems: Principles and Paradigms* (3rd ed.). Pearson Education.
- [2] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal Policy Optimization Algorithms. arXiv preprint arXiv:1707.06347. <https://arxiv.org/abs/1707.06347>
- [3] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- [4] Hinton, G. E., & Salakhutdinov, R. R. (2006). Reducing the dimensionality of data with neural networks. *Science*, 313(5786), 504–507. <https://doi.org/10.1126/science.1127647>
- [5] Liu, B., Xiao, Y., Fu, X., & Yang, H. H. (2021). Trust management in mobile ad hoc networks for secure routing: A review. *Journal of Network and Computer Applications*, 179, 102992. <https://doi.org/10.1016/j.jnca.2021.102992>
- [6] Abadi, M., Barham, P., Chen, J., et al. (2016). TensorFlow: A system for large-scale machine learning. In 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16) (pp. 265–283).
- [7] Paszke, A., Gross, S., Massa, F., et al. (2019). PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems* (pp. 8024–8035).
- [8] Oliviera, R. C., Braun, T., & Heusse, M. (2010). A secure clock synchronization protocol for wireless sensor networks. In *Proceedings of the IEEE Conference on Local Computer Networks (LCN)*, 2010.
- [9] OpenAI. (2020). Gym: A toolkit for developing and comparing reinforcement learning algorithms. GitHub Repository. <https://github.com/openai/gym>
- [10] Lamport, L. (1978). Time, clocks, and the ordering of events in a distributed system. *Communications of the ACM*, 21(7), 558–565. <https://doi.org/10.1145/359545.359563>